



UNIVERSITAT POLITÈCNICA DE CATALUNYA
BARCELONATECH

Departament d'Arquitectura de Computadors

Tarjetas Gráficas y Aceleradores


CUDA – Sesión 02 - Puzzles

Agustín Fernández

Departament d'Arquitectura de Computadors

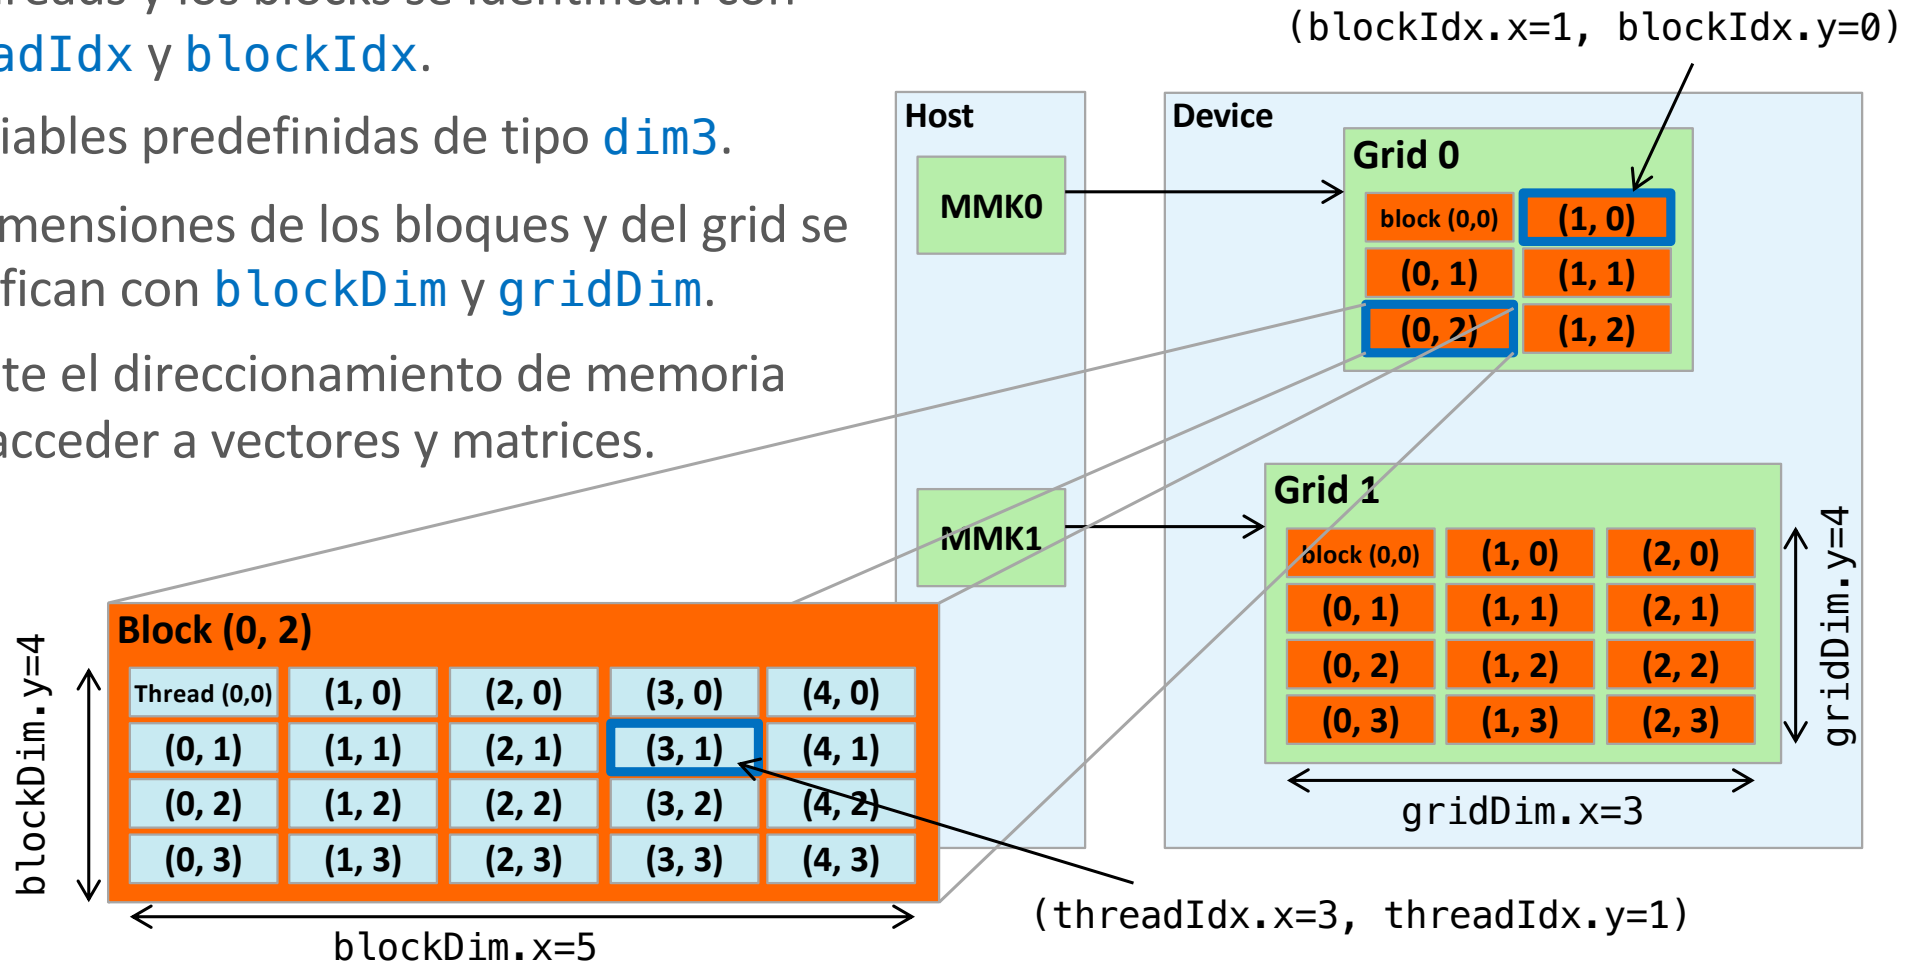
Facultat d'Informàtica de Barcelona

Universitat Politècnica de Catalunya



blockIdx & threadIdx

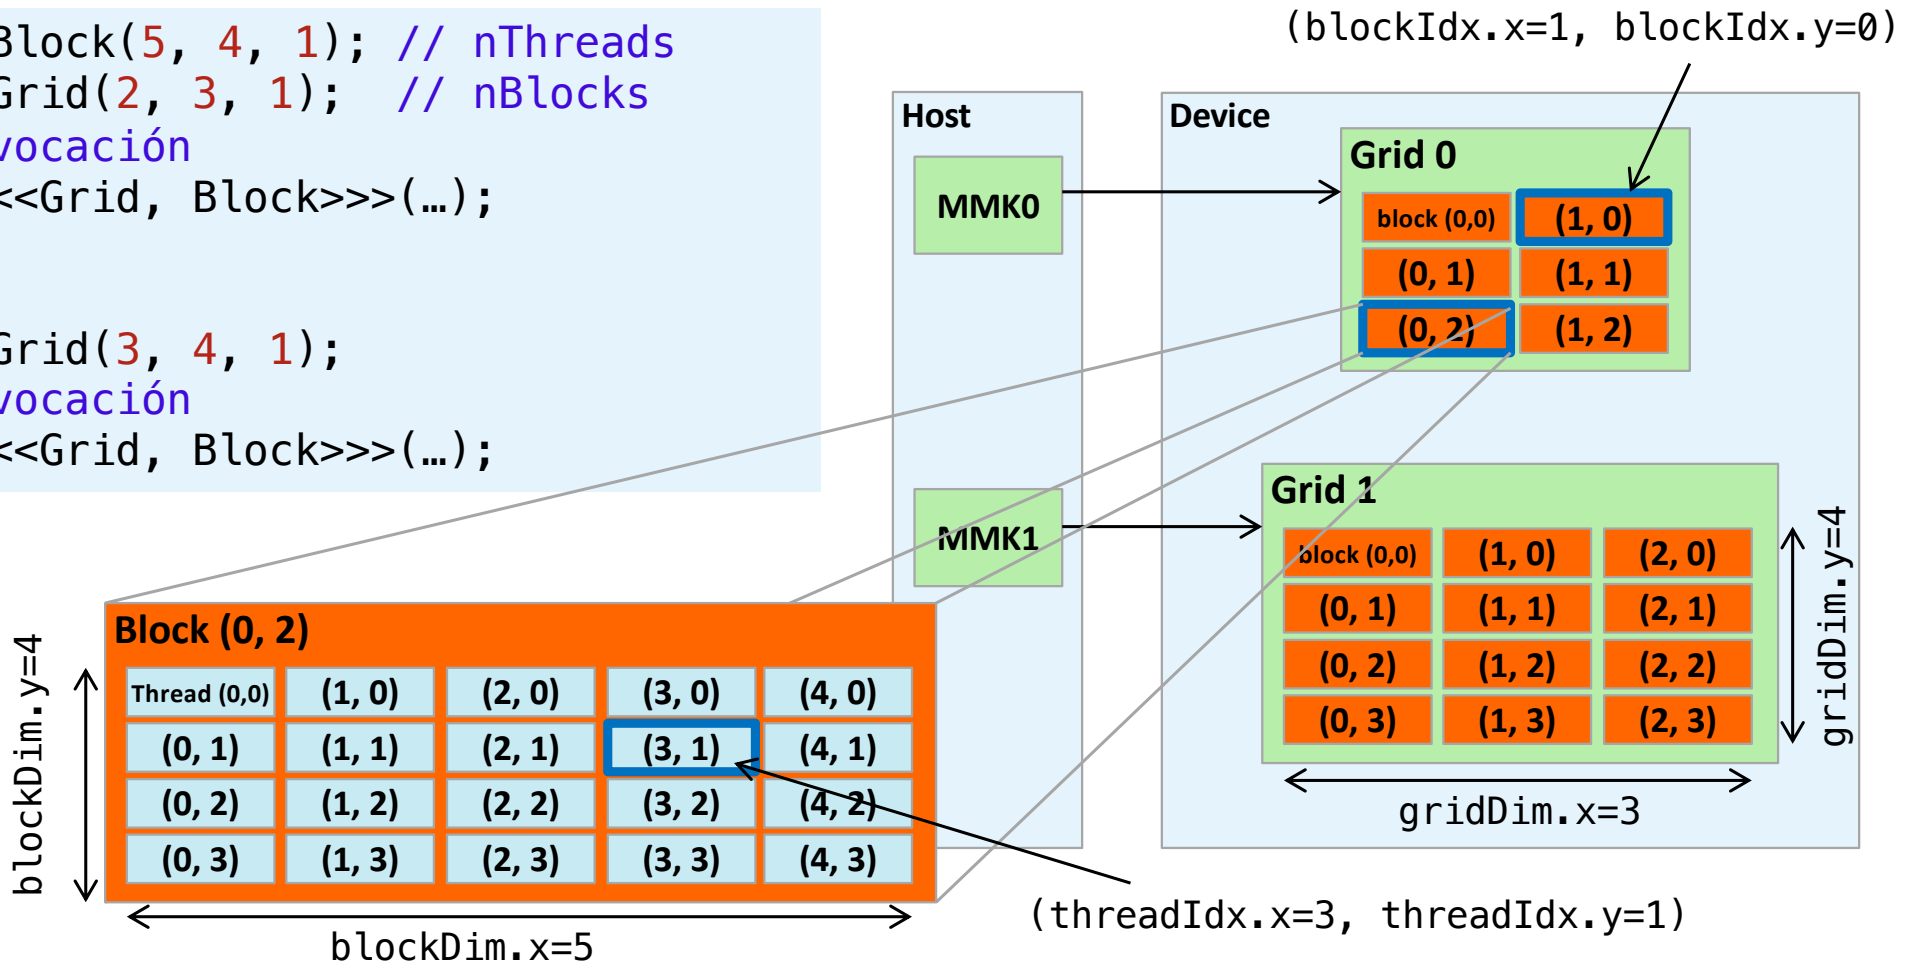
- ❑ Los threads y los blocks se identifican con `threadIdx` y `blockIdx`.
 - Variables predefinidas de tipo `dim3`.
- ❑ Las dimensiones de los bloques y del grid se identifican con `blockDim` y `gridDim`.
- ❑ Permite el direccionamiento de memoria para acceder a vectores y matrices.



blockIdx & threadIdx

```
dim3 Block(5, 4, 1); // nThreads
dim3 Grid(2, 3, 1);  // nBlocks
// Invocación
MMK0<<<Grid, Block>>>(...);
```

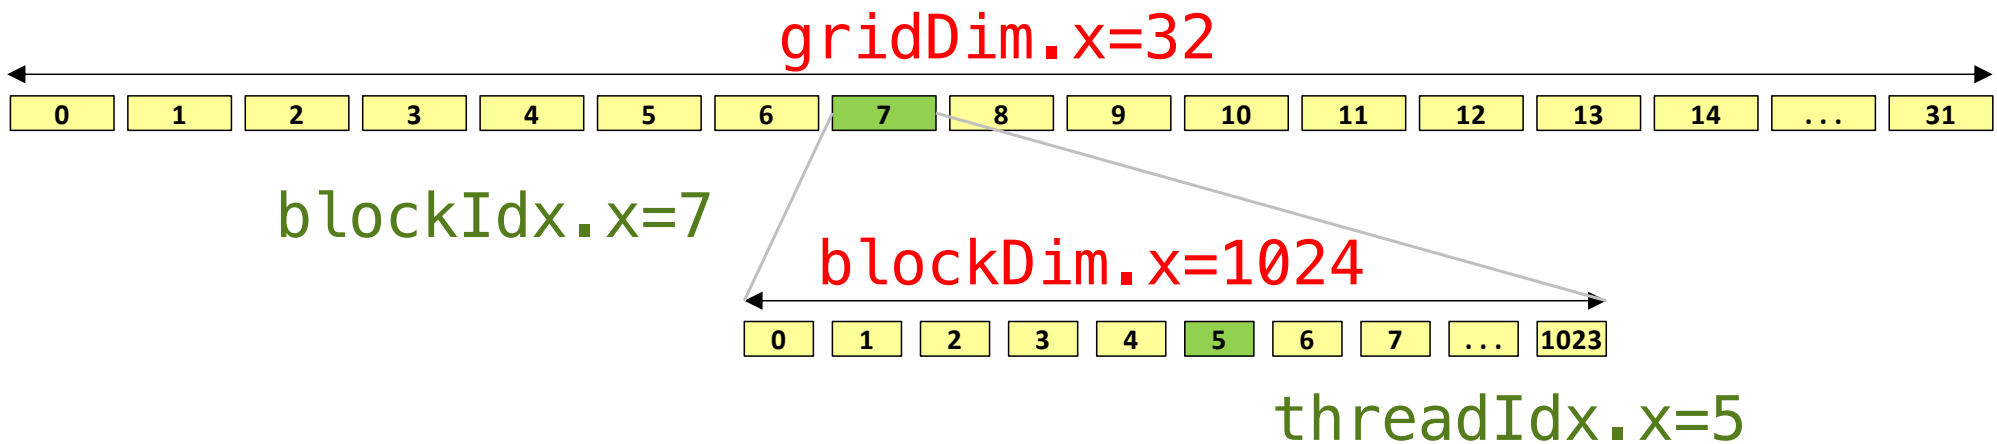
```
dim3 Grid(3, 4, 1);
// Invocación
MMK1<<<Grid, Block>>>(...);
```



Puzzle 1D

$$N = 32 \cdot 1024$$

```
dim3 dimBlock(1024, 1, 1);  
dim3 dimGrid((N+1023)/1024, 1, 1);  
puzzle1DPAR<<dimGrid, dimBlock>>>(N, . . .);
```



$$Id = blockIdx.x * blockDim.x + threadIdx.x = 7 \cdot 1024 + 5 = 7173$$

Puzzle1D

```
void puzzle1DSeq(int N, float *z, float *x, float *y) {  
    int i;  
    for (i=0; i<N; i++)  
        z[i] = 0.5*x[i] + 0.75*y[i] + x[i]*y[i];  
}
```

```
__global__ void puzzle1DPAR(int N, float *z, float *x, float *y) {  
    int i = blockIdx.x * blockDim.x + threadIdx.x;  
    z[i] = 0.5*x[i] + 0.75*y[i] + x[i]*y[i];  
}
```

```
nThreads = 1024;  
nBlocks = N/nThreads;  
dim3 dimGrid(nBlocks, 1, 1);  
dim3 dimBlock(nThreads, 1, 1);  
puzzle1DPAR<<<dimGrid, dimBlock>>>(N, dZ, dX, dY);
```


Puzzle1D. Cuestiones

2. Cambiad el tamaño para que **NO** sea múltiplo del número de threads. Modificad el código, dónde sea necesario, para hacer que funcione correctamente.

```
__global__ void puzzle1DPAR(int N, float *z, float *x, float *y) {  
    int i = blockIdx.x * blockDim.x + threadIdx.x;  
    if (i < N)  
        z[i] = 0.5*x[i] + 0.75*y[i] + x[i]*y[i];  
}  
  
nThreads = 1024;  
nBlocks = (N + nThreads - 1)/nThreads;  
dim3 dimGrid(nBlocks, 1, 1);  
dim3 dimBlock(nThreads, 1, 1);  
puzzle1DPAR<<<dimGrid, dimBlock>>>(N, dZ, dX, dY);
```

Puzzle1D. Cuestiones

3. ¿Qué modificarías en la implementación para considerar ahora que el tamaño del problema **NO** cabe en la memoria de la GPU?




```
cudaMemcpy(..., M, HtD);  
kernel<<<...>>>(M, ...);  
cudaMemcpy(..., M, DtH);
```

Hay que trocear el problema, para que los subproblemas quepan en memoria. Sólo hay que manipular punteros.


Puzzle1D. Cuestiones

4. ¿Qué modificarías en la implementación si el número de Blocks es fijo?



```
__global__ void puzzle1DPAR(int N, float *z, float *x, float *y) {  
    int i = blockIdx.x * blockDim.x + threadIdx.x;  
    int stride = blockDim.x * gridDim.x;  
  
    while (i < N) {  
        z[i] = 0.5*x[i] + 0.75*y[i] + x[i]*y[i];  
        i = i + stride;  
    }  
}  
  
puzzle1DPAR<<10000, 1024>>>(N, dZ, dX, dY);
```


Puzzle 2D por columnas



```
dim3 dimBlockC(8, 1, 1);  
dim3 dimGridC((22+7)/8, 1, 1);  
puzzle2DPAR<<<dimGridC, dimBlockC>>>(N, . . .);
```

$col = blockIdx.x * blockDim.x + threadIdx.x$


Puzzle2D por columnas

```
void puzzle2DSeq(int Nfil, int Ncol, float *z, float *x, float *y) {
    int i, j, ind;
    for (i=0; i<Nfil; i++)
        for (j=0; j<Ncol; j++){
            ind = i * Ncol + j;
            z[ind] = 0.5*x[ind] + 0.75*y[ind] + x[ind]*y[ind];
        }
}
```

```
__global__ void puzzle2DPARcol(int Nfil, int Ncol, float *z, float *x, float *y) {
    int j = blockIdx.x * blockDim.x + threadIdx.x;
    int ind = j;
    if (j < Ncol)
        for (int i=0; i<Nfil; i++, ind = ind + Ncol)
            z[ind] = 0.5*x[ind] + 0.75*y[ind] + x[ind]*y[ind];
}

nThreads = 256;
nBlocks = (Ncol + nThreads - 1)/nThreads;
dim3 dimGridC(nBlocks, 1, 1);
dim3 dimBlockC(nThreads, 1, 1);
puzzle2DPARcol<<<dimGridC, dimBlockC>>>(Nfil, Ncol, dZ, dX, dY);
```

Puzzle 2D por filas



```
dim3 dimBlockF(1, 8, 1);  
dim3 dimGridF(1, (22+7)/8, 1);  
puzzle2DPAR<<<dimGridF, dimBlockF>>>(N, . . .);
```

$$\text{fil} = \text{blockIdx.y} * \text{blockDim.y} + \text{threadIdx.y}$$

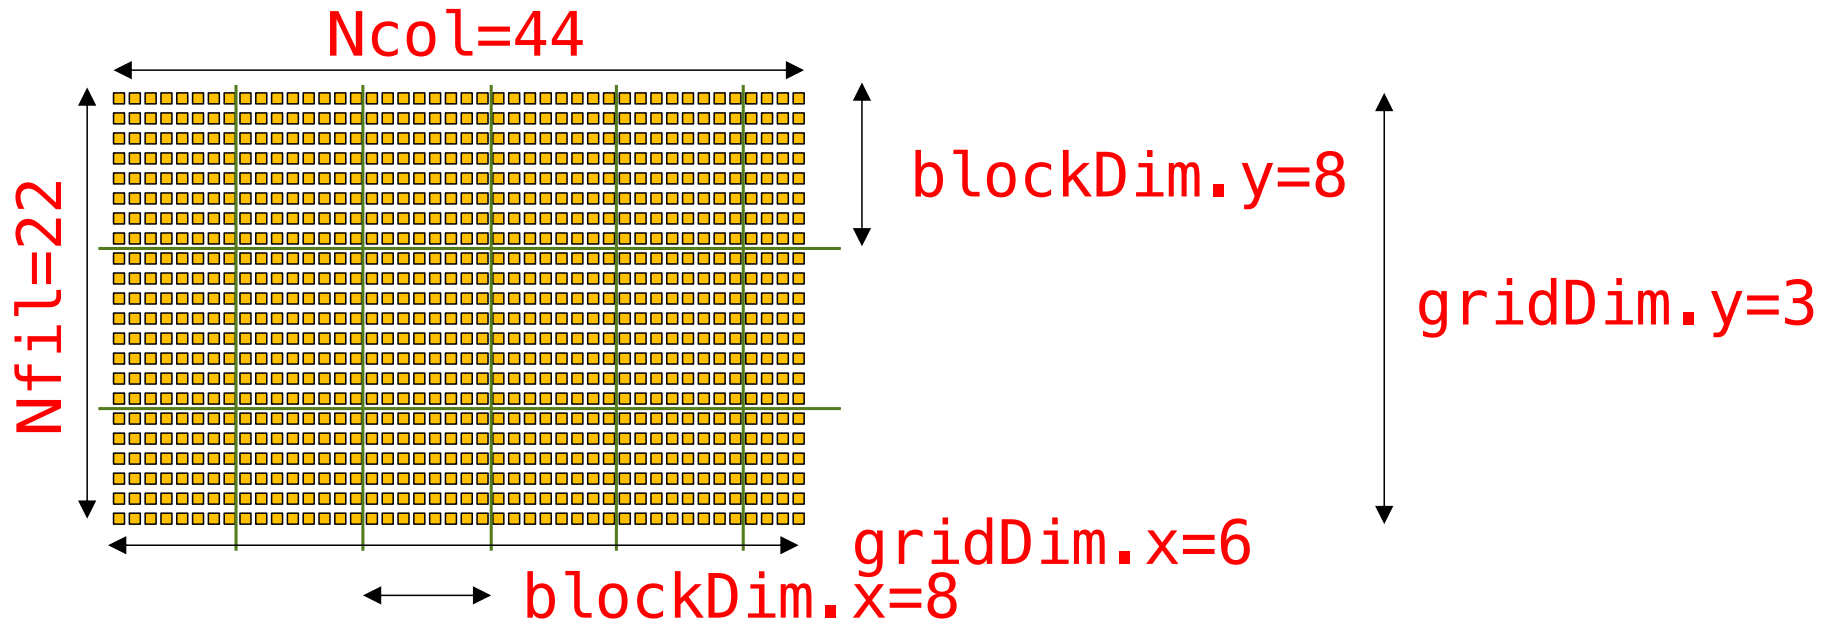
Puzzle2D por filas

```
void puzzle2DSeq(int Nfil, int Ncol, float *z, float *x, float *y) {
    int i, j, ind;
    for (i=0; i<Nfil; i++)
        for (j=0; j<Ncol; j++){
            ind = i * Ncol + j;
            z[ind] = 0.5*x[ind] + 0.75*y[ind] + x[ind]*y[ind];
        }
}
```

```
__global__ void puzzle2DPARfil(int Nfil, int Ncol, float *z, float *x, float *y) {
    int i = blockIdx.y * blockDim.y + threadIdx.y;
    int ind = i * Ncol;
    if (i < Nfil)
        for (int j=0; j<Ncol; j++, ind++)
            z[ind] = 0.5*x[ind] + 0.75*y[ind] + x[ind]*y[ind];
}

nThreads = 256;
nBlocks = (Nfil + nThreads - 1)/nThreads;
dim3 dimGridF(1, nBlocks, 1);
dim3 dimBlockF(1, nThreads, 1);
puzzle2DPARfil<<<dimGridF, dimBlockF>>>(Nfil, Ncol, dZ, dX, dY);
```

Puzzle 2D elemento a elemento



```
dim3 dimBlockE(8, 8, 1);  
dim3 dimGridE((44+7)/8, (22+7)/8, 1);  
puzzle2DPAR<<<dimGridE, dimBlockE>>>(N, . . .);
```

```
fil = blockIdx.y * blockDim.y + threadIdx.y  
col = blockIdx.x * blockDim.x + threadIdx.x
```

Puzzle2D elemento a elemento

```
void puzzle2DSeq(int Nfil, int Ncol, float *z, float *x, float *y) {
    int i, j, ind;
    for (i=0; i<Nfil; i++)
        for (j=0; j<Ncol; j++){
            ind = i * Ncol + j;
            z[ind] = 0.5*x[ind] + 0.75*y[ind] + x[ind]*y[ind];
        }
}
```

```
__global__ void puzzle2DPAR1x1(int Nfil, int Ncol, float *z, float *x, float *y) {
    int i = blockIdx.y * blockDim.y + threadIdx.y;
    int j = blockIdx.x * blockDim.x + threadIdx.x;
    int ind = i * Ncol + j;
    if (i < Nfil && j < Ncol)
        z[ind] = 0.5*x[ind] + 0.75*y[ind] + x[ind]*y[ind];
}

nThreads = 16;
nBlocksCol = (Ncol + nThreads - 1)/nThreads;
nBlocksFil = (Nfil + nThreads - 1)/nThreads;
dim3 dimGridE(nBlocksCol, nBlocksFil, 1);
dim3 dimBlockE(nThreads, nThreads, 1);
puzzle2DPAR1x1<<<dimGridE, dimBlockE>>>(Nfil, Ncol, dZ, dX, dY);
```

Puzzle2D columnas vs filas vs elementos

```
__global__ void puzzle2DPARcol(int Nfil, int Ncol, float *z, float *x, float *y) {  
    int j = blockIdx.x * blockDim.x + threadIdx.x;  
    int ind = j;  
    if (j < Ncol)  
        for (int i=0; i<Nfil; i++, ind = ind + Ncol)  
            z[ind] = 0.5*x[ind] + 0.75*y[ind] + x[ind]*y[ind];  
}
```

```
nTh = 256;  
nBlC = (Ncol + nTh-1)/nTh;  
dim3 dimGridC(nBlC, 1, 1);  
dim3 dimBlockC(nTh, 1, 1);
```

```
__global__ void puzzle2DPARfil(int Nfil, int Ncol, float *z, float *x, float *y) {  
    int i = blockIdx.y * blockDim.y + threadIdx.y;  
    int ind = i * Ncol;  
    if (i < Nfil)  
        for (int j=0; j<Ncol; j++, ind++)  
            z[ind] = 0.5*x[ind] + 0.75*y[ind] + x[ind]*y[ind];  
}
```

```
nTh = 256;  
nBlF = (Nfil + nTh-1)/nTh;  
dim3 dimGridF(1, nBlF, 1);  
dim3 dimBlockF(1, nTh, 1);
```

```
__global__ void puzzle2DPAR1x1(int Nfil, int Ncol, float *z, float *x, float *y) {  
    int i = blockIdx.y * blockDim.y + threadIdx.y;  
    int j = blockIdx.x * blockDim.x + threadIdx.x;  
    int ind = i * Ncol + j;  
    if (i < Nfil && j < Ncol)  
        z[ind] = 0.5*x[ind] + 0.75*y[ind] + x[ind]*y[ind];  
}
```

```
nTh = 16;  
nBlC = (Ncol + nTh-1)/nTh;  
nBlF = (Nfil + nTh-1)/nTh;  
dim3 dimGridE(nBlC, nBlF, 1);  
dim3 dimBlockE(nTh, nTh, 1);
```

Puzzle2D columnas vs filas vs elementos

```
nThreads = 256;
nBlocks = (Ncol + nThreads - 1)/nThreads;
dim3 dimGridC(nBlocks, 1, 1);
dim3 dimBlockC(nThreads, 1, 1);
puzzle2DPARcol<<<dimGridC, dimBlockC>>>(Nfil, Ncol, dZ, dX, dY);
```

```
nThreads = 256;
nBlocks = (Nfil + nThreads - 1)/nThreads;
dim3 dimGridF(1, nBlocks, 1);
dim3 dimBlockF(1, nThreads, 1);
puzzle2DPARfil<<<dimGridF, dimBlockF>>>(Nfil, Ncol, dZ, dX, dY);
```

```
nThreads = 16;
nBlocksCol = (Ncol + nThreads - 1)/nThreads;
nBlocksFil = (Nfil + nThreads - 1)/nThreads;
dim3 dimGridE(nBlocksCol, nBlocksFil, 1);
dim3 dimBlockE(nThreads, nThreads, 1);
puzzle2DPAR1x1<<<dimGridE, dimBlockE>>>(Nfil, Ncol, dZ, dX, dY);
```


Puzzle2D Rendimiento

Kernel por Filas

Dimension problema: 1023 filas x 1023 columnas

Dimension Block: 1 x 256 x 1 (256) threads

Dimension Grid: 1 x 4 x 1 (4) blocks

Kernel por Columnas

Dimension problema: 1023 filas x 1023 columnas

Dimension Block: 256 x 1 x 1 (256) threads

Dimension Grid: 4 x 1 x 1 (4) blocks

Kernel Elemento a Elemento

Dimension problema: 1023 filas x 1023 columnas

Dimension Block: 16 x 16 x 1 (256) threads

Dimension Grid: 64 x 64 x 1 (4096) blocks

Resumen Rendimiento

Tiempo Paralelo Kernel filas: 0.670912 ms (7.80 GFLOPS)

Tiempo Paralelo Kernel columnas: 0.640416 ms (8.17 GFLOPS)

Tiempo Paralelo Kernel elemento a elemento: 0.052288 ms (100.07 GFLOPS)

Tiempo Secuencial: 1.088600 milseg (4.81 GFLOPS)

¿Explicación de los rendimientos?


Número de Bloques

Patrones de acceso a memoria de los warps

RTX 3090


Puzzle2D Rendimiento

¿Explicación de los rendimientos?




Kernel por Filas
4 blocks (256 threads)
7.80 GFLOPS (RTX 3080)
1.92 GFLOPS (K40c)
6.97 GFLOPS (GTX 1080Ti)

Patrones de acceso a memoria de los warps



Kernel por Columnas
4 blocks (256 threads)
8.17 GFLOPS (RTX 3080)
7.09 GFLOPS (K40c)
9.62 GFLOPS (GTX 1080Ti)



Kernel Elemento a Elemento
4096 blocks (16x16 threads)
100.07 GFLOPS (RTX 3080)
44.59 GFLOPS (K40c)
78.77 GFLOPS (GTX 1080Ti)

Puzzle3D elemento a elemento

```
void puzzle3DSeq(int Ncar, int Nfil, int Ncol, float *z, float *x, float *y) {  
    int i, j, t, ind;  
    for (t=0; t<Ncar; t++)  
        for (i=0; i<Nfil; i++)  
            for (j=0; j<Ncol; j++) {  
                ind = t*Nfil*Ncol + i*Ncol + j;  
                z[ind] = 0.5*x[ind] + 0.75*y[ind] + x[ind]*y[ind];  
            }  
}
```

```
__global__ void puzzle3DPAR1x1x1(int Ncar, int Nfil, int Ncol, float *z, float *x, float *y){  
    int i = blockIdx.y * blockDim.y + threadIdx.y;  
    int j = blockIdx.x * blockDim.x + threadIdx.x;  
    int t = blockIdx.z * blockDim.z + threadIdx.z;  
    int ind = t*Nfil*Ncol + i*Ncol + j;  
    if (t<Ncar && i<Nfil && j<Ncol)  
        z[ind] = 0.5*x[ind] + 0.75*y[ind] + x[ind]*y[ind];  
}
```

Puzzle3D elemento a elemento

```
int nThreads = 8;
int nBlocksFil = (Nfil+nThreads-1)/(nThreads);
int nBlocksCol = (Ncol+nThreads-1)/nThreads;
int nBlocksCar = (Ncar+nThreads-1)/nThreads;

dim3 dimGridE(nBlocksCol, nBlocksFil, nBlocksCar);
dim3 dimBlockE(nThreads, nThreads, nThreads);
puzzle3DPAR1x1x1<<<dimGridE, dimBlockE>>>(Ncar, Nfil, Ncol, dZ, dX, dY);
```

```
__global__ void puzzle3DPAR1x1x1(int Ncar, int Nfil, int Ncol, float *z, float *x, float *y){
    int i = blockIdx.y * blockDim.y + threadIdx.y;
    int j = blockIdx.x * blockDim.x + threadIdx.x;
    int t = blockIdx.z * blockDim.z + threadIdx.z;
    int ind = t*Nfil*Ncol + i*Ncol + j;
    if (t<Ncar && i<Nfil && j<Ncol)
        z[ind] = 0.5*x[ind] + 0.75*y[ind] + x[ind]*y[ind];
}
```

Uso de nvprof

```
nsys nvprof --print-gpu-summary ./puzzle2D.exe 1024 1024 Y
```

CUDA Kernel Statistics:

Time(%)	Total Time (ns)	Instances	Average	Minimum	Maximum	Name
49.2	643,428	1	643,428.0	643,428	643,428	puzzle2DPARfil(int, int, float*, float*, float*)
47.8	624,741	1	624,741.0	624,741	624,741	puzzle2DPARcol(int, int, float*, float*, float*)
3.0	39,424	1	39,424.0	39,424	39,424	puzzle2DPAR1x1(int, int, float*, float*, float*)

CUDA Memory Operation Statistics (by time):

Time(%)	Total Time (ns)	Operations	Average	Minimum	Maximum	Operation
65.2	977,575	3	325,858.3	325,730	325,923	[CUDA memcpy DtoH]
34.8	521,284	2	260,642.0	259,041	262,243	[CUDA memcpy HtoD]

CUDA Memory Operation Statistics (by size in KiB):

Total	Operations	Average	Minimum	Maximum	Operation
8,176.008	2	4,088.004	4,088.004	4,088.004	[CUDA memcpy HtoD]
12,264.012	3	4,088.004	4,088.004	4,088.004	[CUDA memcpy DtoH]

También da una lista de todas las llamadas a la API de CUDA realizadas

Uso de nvprof

```
nsys nvprof --print-gpu-trace ./puzzle2D.exe 1024 1024 Y
```

CUDA Kernel & Memory Operations Trace:

Start(sec)	Time(ns)	CorrId	Grid Size	Block Size	Reg/T	StcSMem	DymSMem	Bytes	(GB/s)	SrcMemKd	DstMemKd	Device	Str	Name
0.360799	262,243	426						4,186,116	15.962	Pinned	Device	RTX 3080 (3)	7	[CUDA memcpy HtoD]
0.361076	259,041	427						4,186,116	16.160	Pinned	Device	RTX 3080 (3)	7	[CUDA memcpy HtoD]
0.361421	643,428	431	(1 4 1)	(1 256 1)	24	0	0					RTX 3080 (3)	7	puzzle2DPAFil()
0.362078	325,730	435						4,186,116	12.851	Device	Pinned	RTX 3080 (3)	7	[CUDA memcpy DtoH]
0.364149	624,741	439	(4 1 1)	(256 1 1)	24	0	0					RTX 3080 (3)	7	puzzle2DPAcol()
0.364789	325,922	443						4,186,116	12.843	Device	Pinned	RTX 3080 (3)	7	[CUDA memcpy DtoH]
0.366808	39,424	447	(64 64 1)	(16 16 1)	18	0	0					RTX 3080 (3)	7	puzzle2DPA1x1()
0.366861	325,923	451						4,186,116	12.843	Device	Pinned	RTX 3080 (3)	7	[CUDA memcpy DtoH]

Usando todas las GPUs

En Boada-10 hay 4 GPUs. Si no hacemos nada, TODOS los jobs se ejecutan en la GPU 0.

En el `job.sh`

```
...  
#SBATCH --gres=gpu:4  
...
```

```
cudaGetDeviceCount(&count);  
srand(time(NULL));  
gpu = rand();  
cudaSetDevice((gpu>>3) % count);
```

La rutina `cudaGetDeviceCount` nos dice cuantas GPUs hay en el sistema, el resultado estará en `count`.

Obtiene un número aleatorio. Nos aseguramos que sea un valor diferente cada vez que ejecutamos el programa.

La rutina `cudaSetDevice(X)` indica que todo lo que viene a continuación se ejecuta en la GPU `X`. `X` ha de ser un valor entre `0` y `count-1`.



UNIVERSITAT POLITÈCNICA DE CATALUNYA
BARCELONATECH

Departament d'Arquitectura de Computadors

Tarjetas Gráficas y Aceleradores

CUDA – Sesión 02 - Puzzles

Agustín Fernández

Departament d'Arquitectura de Computadors

Facultat d'Informàtica de Barcelona

Universitat Politècnica de Catalunya

